

FINAL 14-DAY LLM DEVELOPER ROADMAP

Interview-Focused Plan | LLM Developer Preparation

Filtered from a 33-hour RAG Bootcamp to the highest ROI topics for a junior/fresher LLM Developer interview.

Day(s)	Phase	Hours
1	RAG Fundamentals	1.5
2-3	Embeddings & Vector DBs	2.5
4-5	Advanced RAG	2
5-6	LangGraph & Agents	4
7-8	Agentic RAG	4.5
9	Evaluation	1
10	KV Cache / vLLM	0.7
11-12	Projects	1.5
13	Transformers / SQL / Pandas / Git	4
14	Revision + Mock Interview	-

Phase 1: RAG Fundamentals

- Introduction to RAG
- Advantages of RAG
- Prompt Engineering vs Fine-Tuning vs RAG
- Data Ingestion & Preprocessing
- Query Processing
- Text Splitting Techniques

Phase 2: Embeddings & Vector Databases

- Embeddings
- Cosine Similarity
- HuggingFace Embeddings
- Semantic Search
- Vector Stores vs Vector Databases
- Traditional RAG with ChromaDB

Phase 3: Advanced RAG

- Conversational Memory
- Semantic Chunking
- Query Expansion
- Query Decomposition
- HyDE

Phase 4: LangGraph & Agents

- AI Agents vs Agentic AI
- Tools in LangChain
- LangGraph Basics
- Router & Tool Integration

Phase 5: Agentic RAG

- ReAct Agents
- Agentic RAG
- Multi-Agent RAG
- Corrective RAG
- Adaptive RAG

Phase 6: Evaluation

- LLM-as-a-Judge
- RAG Evaluation
- Metrics

- Evaluator Pipelines

Phase 7: Production Topics

- KV Cache
- PagedAttention
- vLLM
- Inference Parameters

Phase 8: Portfolio Projects

- Document Ingestion
- Vector Store
- Retriever
- Graph Builder
- ReAct Agent
- Streamlit App

Phase 9: Extras

- Transformers Basics
- LoRA / QLoRA
- Pandas
- SQL
- Git & GitHub

Interview Questions Checklist

- Explain your RAG pipeline architecture.
- What is semantic chunking and why use it?
- What are embeddings?
- Vector Store vs Vector Database?
- LangChain vs LangGraph?
- What is a ReAct Agent?
- How do you evaluate a RAG system?
- What is KV Cache?
- What is PagedAttention?
- How would you optimize inference latency?

Expected Outcome

By the end of this roadmap you should be able to build and explain a RAG application, use LangChain and LangGraph, implement Agentic RAG workflows, discuss evaluation strategies, and confidently answer common junior LLM Developer interview questions.